

VIVEK SURYAVANSHI

Senior GenAI Engineer | Solutions Architect | 4x Anthropic Claude Certified

New Delhi, India | vivek.suryavanshi@gmail.com | +91-9560613574 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Senior GenAI engineer and solutions builder with 14+ years in enterprise software delivery, 4+ years building agentic AI, RAG, and LLM systems in production. I take AI systems from POC to production — designing agentic workflows with LangGraph and agent SDKs, integrating enterprise tools and data, and building rigorous evals across agent layers to assess quality. Anthropic Claude-certified across Architect and Developer tracks.

CORE TECHNICAL SKILLS

Agentic AI & Orchestration: LangGraph, OpenAI Agents SDK, LangChain, CrewAI, planner-reasoner & ReAct workflows, memory & state management, tool orchestration & function calling, agentic RAG, guardrails, prompt engineering (CoT, few-shot)

Agentic Development & Tooling: Claude Code (skills, hooks, CLAUDE.md configuration), Model Context Protocol (MCP) server implementation & integration (with tools), agent harness design, spec-driven development

Agentic & LLM Evaluation: Multi-level eval design (end-to-end, trajectory-level, component-level), golden dataset construction, LLM-as-judge scoring (faithfulness, retrieval accuracy, context relevance), tool-selection & parameter-accuracy benchmarking, Ragas, eval-driven quality gates for agentic pipelines

LLM Platforms & Models: OpenAI API (GPT-4o, Whisper, TTS), Azure OpenAI, Anthropic Claude (Opus, Sonnet, Haiku), Google Gemini & Vertex AI, Hugging Face Transformers, Cohere, local model deployment (Ollama, MLX)

RAG & Retrieval: RAG architectures, hybrid retrieval (keyword + vector search), multi-source retrieval, chunking & embeddings, enterprise metadata strategies, Azure AI Search, FAISS, Pinecone

LLMOps & Observability: Langfuse — tracing, cost & latency monitoring

Cloud & Backend: Microsoft Azure (Azure OpenAI, AI Search, Cosmos DB, Blob Storage); GCP (Vertex AI, Gemini); Python, FastAPI, Pydantic, REST APIs, Streamlit, Supabase; Docker, Render, Git, Jenkins CI/CD

Delivery & Leadership: AI solution architecture, AI roadmap planning, technical discovery sessions, KPI definition, stakeholder alignment, workshop facilitation, use-case prioritization

PROFESSIONAL EXPERIENCE

Senior GenAI/LLM Engineer / Solutions Architect

Mar 2022 – Present

Accenture Solutions Pvt Ltd, Gurugram

Agentic AI & Production Deployment

- Architected end-to-end enterprise RAG platforms on Azure, integrating Azure OpenAI, Azure AI Search, Cosmos DB, and Blob Storage — enabling 500+ employees across the US to access policy-grounded answers via an internal AI assistant.
- Designed and implemented agentic AI workflows using LangGraph and LangChain with OpenAI API for LLM inference, applying planner-reasoner and ReAct workflow patterns to orchestrate retrieval, tool execution, and state management across enterprise GenAI use cases.
- Built RAG evaluation and observability pipelines using LLM-as-judge scoring to measure answer faithfulness, retrieval accuracy, and context relevance, with Langfuse tracing enabling iterative quality improvement in production.
- Defined agent evaluation metrics — tool-selection precision, parameter accuracy, and trajectory coverage — using golden datasets to benchmark planner-reasoner workflows, retrieval paths, and tool-use sequences.
- Optimized chunking, embeddings, metadata, and hybrid retrieval (BM25 + vector) — reducing token consumption by ~25% and improving retrieval accuracy across enterprise documents.
- Built FastAPI inference APIs with structured Pydantic outputs, authentication, and backend orchestration — delivering production-grade GenAI inference services.
- Applied prompt engineering techniques including Chain-of-Thought, ReAct, few-shot prompting, and guardrails to improve reasoning accuracy and mitigate hallucinations in production systems.
- Configured and implemented Model Context Protocol (MCP) servers for PostgreSQL and Playwright in VS Code — enabling the QA team to execute agentic test workflows through GitHub Copilot agent mode for database validation and browser-based UI testing.

Technical Discovery & Customer Delivery

- Led 10+ technical discovery and design sessions with cross-functional stakeholders (operations, SMEs, support) — translating complex business problems into actionable GenAI solution blueprints and accelerating sign-off and adoption.
- Co-owned the roadmap for internal GenAI assistants across 3+ initiatives — prioritizing high-value use cases, defining success metrics, and aligning delivery with business leadership.

Business Analyst – Digital & AI Solutions

Mar 2021 – Mar 2022

Accenture Solutions Pvt Ltd, Gurugram

- Translated business requirements into functional specs, KPIs, and success metrics for AI-enabled workflows; ran stakeholder workshops, UAT cycles, and feedback loops driving measurable adoption.

NLP Engineer (Python)

Mar 2019 – Feb 2021

Accenture Solutions Pvt Ltd, Gurugram

- Built Python NLP preprocessing pipelines (spaCy, NLTK) for cleaning, NER, tokenization, and lemmatization; produced structured, entity-enriched JSON consumed by search, analytics, and automation systems — the foundation for later LLM work.

Senior Automation QA Engineer

Apr 2011 – Feb 2019

Accenture Solutions Pvt Ltd, Gurugram

- Led test automation (Selenium, TestNG, Java, SQL) with 75%+ coverage; integrated Jenkins CI/CD pipelines, reducing regression effort by ~40%.

NOTABLE PROJECTS

WhisperRAG – Voice-Powered RAG with Production Observability [Live App](#) | [GitHub](#)

Voice RAG grounded strictly in a loaded webpage. FastAPI + OpenAI Whisper/GPT-4o-mini/TTS, Supabase session logging. Langfuse tracks cost, latency, and traces; LLM-as-judge runs faithfulness scoring on every conversation — demonstrating production-style evaluation + observability.

Cognify – Stateful Multi-Agent Executive Function Assistant [Live App](#) | [GitHub](#)

Stateful multi-agent system converting high-level goals into structured task plans via contextual reasoning and tool orchestration.

GreenMe – Sustainability GenAI Application [Live App](#) | [GitHub](#)

GenAI app tracking carbon footprint with personalized sustainability insights, built with LangChain and Cohere.

ExcelSensei – Natural Language Interface for Excel [Live App](#) | [GitHub](#)

LangChain DataFrame agent enabling natural-language querying and analysis of Excel data — no formulas required.

CERTIFICATIONS



- Anthropic Claude Certified Architect – Professional (CCAR-P)
- Anthropic Claude Certified Architect – Foundations (CCAR-F)
- Anthropic Claude Certified Associate – Foundations (CCAO-F)
- Anthropic Claude Certified Developer – Foundations (CCDV-F)
- Microsoft Azure AI Engineer Associate (AI-102)
- Google Generative AI Leader
- AWS Certified Cloud Practitioner (AWS - CCP)
- Databricks Certified Generative AI Engineer – Associate
- BCS ISEB Certified Business Analyst

EDUCATION

B.Tech – Electronics & Communication Engineering (MDU, Rohtak)

2006 – 2010